

The RBM mutual-information lower bound is fixed by the marginals, not by a learning principle

Vicente Humberto Monteverde

Universidad del Museo Social Argentino (UMSA), Argentina

ORCID 0000-0001-8884-4811

Draft of 27 September 2026

Abstract

A recent review [1] (Sec. 3.2 and Fig. 4 there, after Ref. [34] of the review) characterizes a restricted Boltzmann machine (RBM) learner by two numbers for each visible–hidden pair. The first is the covariance $\eta = \text{Cov}(\sigma_z^v, \sigma_z^h)$, which is read from the imaginary part of an out-of-time-order correlator. The second is the mutual information I . The pair (I, η) is shown to lie between analytic bounds $\text{LB}(\eta)$ and $\text{UB}(\eta)$. Trained RBMs sit on the lower bound for every system size and field ratio of transverse-field Ising drivers, and the review reads this as a learning principle: the network would use the least mutual information compatible with the covariance. We point out that a pair of ± 1 units is fixed by three numbers, the two marginals $m_v = \langle \sigma_z^v \rangle$, $m_h = \langle \sigma_z^h \rangle$ and η . In those variables, $\text{LB}(\eta)$ is exactly the case $m_v = m_h = 0$, and it is the minimum of I over the marginals at fixed η . $\text{UB}(\eta)$ is the perfectly correlated pair with $|m_v| = |m_h| = \sqrt{1 - |\eta|}$. The transverse-field Ising driver is $\mathbb{Z}_2$ -symmetric, so a symmetric RBM has zero marginals and sits on LB by construction. We test the consequence that could have failed. With exact enumeration ($n = 6$ visible spins, $\alpha = 1$), a longitudinal field of 0.1–0.3 moves the same learner off LB by up to 0.17 bits, in step with the measured marginals. In the ordered phase, a symmetry-broken and a symmetric RBM of the same target sit off and on LB respectively. The OTOC reading of η is unaffected; we suggest reporting the marginals alongside (I, η) .

1 The bounds in terms of the marginals

For two ± 1 variables s (visible) and t (hidden) with marginals m_v, m_h and covariance $\eta = \langle st \rangle - m_v m_h$, the joint distribution is

$$p(s, t) = \frac{1}{4} [1 + s m_v + t m_h + st (\eta + m_v m_h)], \quad (1)$$

and the mutual information (in bits) is

$$I(m_v, m_h, \eta) = h\left(\frac{1+m_v}{2}\right) + h\left(\frac{1+m_h}{2}\right) - H(p), \quad (2)$$

with h the binary entropy and $H(p)$ the entropy of Eq. (1). With $\ell(x) = -x \log_2 x$, the bounds of Ref. [1] read $\text{LB}(\eta) = 2 - 2\ell(\frac{1+\eta}{4}) - 2\ell(\frac{1-\eta}{4})$ and $\text{UB}(\eta) = \ell(\frac{1}{2} + \frac{1}{2}\sqrt{1-|\eta|}) + \ell(\frac{1}{2} - \frac{1}{2}\sqrt{1-|\eta|})$. Base-2 logarithms reproduce the plotted endpoints $\text{LB}(0) = 0$ and $\text{LB}(1) = \text{UB}(1) = 1$.

Proposition 1. (i) $\text{LB}(\eta) = I(0, 0, \eta)$. (ii) $\text{UB}(\eta) = I(r, r, \eta)$ with $r = \sqrt{1 - |\eta|}$ (sign of m_h following that of η), which is the locked pair $t = \pm s$. (iii) To leading order in η and the marginals,

$$I - \text{LB}(\eta) \simeq \frac{\eta^2 (m_v^2 + m_h^2)}{2 \ln 2}. \quad (3)$$

Table 1: Trained RBMs (median over 3 seeds, 36 pairs each). Fidelity is with the exact ground state; gaps are in bits.

$g = B/J$	h_z	fidelity	$\langle m_v \rangle$	mean $I - \text{LB}$	max $I - \text{LB}$
1.0	0	1.0000	0.0001	3×10^{-7}	1×10^{-5}
2.0	0	1.0000	0.0000	7×10^{-7}	2×10^{-5}
1.0	0.1	1.0000	0.39	1.6×10^{-2}	0.10
1.0	0.3	1.0000	0.52	1.4×10^{-2}	0.17
2.0	0.1	1.0000	0.07	5.2×10^{-4}	5×10^{-3}
2.0	0.3	1.0000	0.18	2.7×10^{-3}	3×10^{-2}
0.5	0 (random start)	0.50	0.81	1.1×10^{-2}	0.12
0.5	0 ($a = b = 0$ start)	1.0000	0.0002	1.5×10^{-8}	1×10^{-7}

(i) and (ii) follow by substitution into Eq. (1). For (iii), expand I in the correlation coefficient $\rho = \eta / \sqrt{(1 - m_v^2)(1 - m_h^2)}$, using $I \simeq \rho^2 / (2 \ln 2)$. Numerically, on a grid of marginals, $\text{LB}(\eta)$ is the minimum of I at fixed η and $\text{UB}(\eta)$ its supremum. Eq. (3) is 5% low at $\eta = 0.1$ and 12% low at $\eta = 0.3$. The whole (I, η) region is therefore swept by the marginals: at fixed covariance, biased units carry more mutual information because their variances are smaller (Fig. 1, left).

2 Why trained RBMs sit on the lower bound

The learner’s state is the thermal distribution of $H(X) = \sum a_i s_i + \sum b_j t_j + \sum W_{ij} s_i t_j$. The driver $H_{\text{drv}} = -B \sum_i X_i - J \sum_{\langle ij \rangle} Z_i Z_j$ commutes with $\prod_i X_i$, so its ground state has $\langle Z_i \rangle = 0$. An RBM with $a = b = 0$ is invariant under $(s, t) \rightarrow (-s, -t)$, so all its visible and hidden marginals vanish, and by (i) every pair lies exactly on $\text{LB}(\eta)$. Saturation of the lower bound is therefore what $\mathbb{Z}_2$ symmetry implies. It is not evidence that training minimizes mutual information.

This reading makes a prediction that can fail. Break the symmetry, and the points must leave LB by the amount Eq. (3) sets from the measured marginals.

3 Test

We trained RBM wavefunctions $\psi(v) = e^{a \cdot v} \prod_j 2 \cosh(b_j + W_j \cdot v)$ [2] on the driver with an added longitudinal field $-h_z \sum_i Z_i$. The setup: a periodic chain of $n = 6$, $\alpha = 1$, real positive amplitudes (sufficient for this stoquastic driver), exact energy by full enumeration, L-BFGS, and 3 seeds. For every visible–hidden pair of $P(v, h) \propto e^{a \cdot v + b \cdot h + v^\top W h}$ we computed (m_v, m_h, η, I) exactly (Table 1, Fig. 1, right).

- *Symmetric driver.* The review’s observation is reproduced: trained RBMs sit on LB, with gaps below 2×10^{-5} bits, because their marginals vanish.
- *Broken symmetry.* A field of 0.1–0.3 moves the same learner off LB by 5×10^{-4} to 2×10^{-2} bits on average, and up to 0.17 bits for a pair. At $g = 2$, where the marginals are small, Eq. (3) matches the mean gap (5.3×10^{-4} vs 5.2×10^{-4} , and 2.6×10^{-3} vs 2.7×10^{-3}). It underestimates the gap by up to $2\times$ when the marginals are large.
- *Same physics, two positions.* In the ordered phase ($g = 0.5$) a random start breaks the symmetry: the fidelity with the symmetric ground state is 0.50 and the marginals are 0.8, and the points leave LB. A start from $a = b = 0$ keeps the symmetry, finds the symmetric ground state exactly, and the points sit on LB. The target is the same; the position in the (I, η) plane is set by the marginals.




Figure 1: Left: the (I, η) region with the bounds of Ref. [1]; grey curves have fixed marginals $m_v = m_h$. Right: all visible–hidden pairs of trained RBMs at $g = 1$ for three longitudinal fields.

4 Discussion

The OTOC route to η in Ref. [1] and the (I, η) picture remain useful. Our point concerns the interpretation. Lower-bound saturation is the zero-marginal case, and it is fixed by the symmetry of the driver, or broken by the learner or by a field. It does not require a minimum-information principle. A simple way to separate the two readings in larger simulations is to report the marginals $\langle \sigma_z^v \rangle$ and $\langle \sigma_z^h \rangle$ alongside (I, η) , or equivalently the distance to LB predicted by Eq. (3). A learner that minimized mutual information would stay on LB even with $h_z \neq 0$; the RBMs here do not.

Limitations. Our test uses small n with exact enumeration and real RBMs. The stochastic-reconfiguration training with Monte Carlo sampling at larger N used in the review, and its long-range driver, are not reproduced. The claim is limited to the one above.

Code availability. `examples/rbm_mutual_information_qg.py` and its regression tests at <https://github.com/Viny2030/qang> (MIT license), Sec. 47 of `RESEARCH_NOTES.md`. In the notation of that repository, m_v and m_h are the polar biases qg_z of the two units, and $h((1 + qg_z)/2)$ is their qg_S [3].

Acknowledgments and use of AI tools

Generative AI tools (Anthropic Claude Opus 5.5) assisted with code, numerical experiments and drafting. The author framed the question, reviewed and validated all output, including every number reported.

References

- [1] Vinit Singh, Amandeep Singh Bhatia, Mandeep Kaur Saggi, Manas Sajjan, and Sabre Kais. Quantum machine learning for complex systems: paradigms, applications, and challenges. *Academia Quantum*, 3, 2026. doi: 10.20935/AcadQuant8243.
- [2] Giuseppe Carleo and Matthias Troyer. Solving the quantum many-body problem with artificial neural networks. *Science*, 355(6325):602–606, 2017. doi: 10.1126/science.aag2302.

- [3] Vicente Humberto Monteverde. The Qang (qg): A unified angular-probability unit and metric for parametric quantum circuit design. 2026. doi: 10.5281/zenodo.22832150.